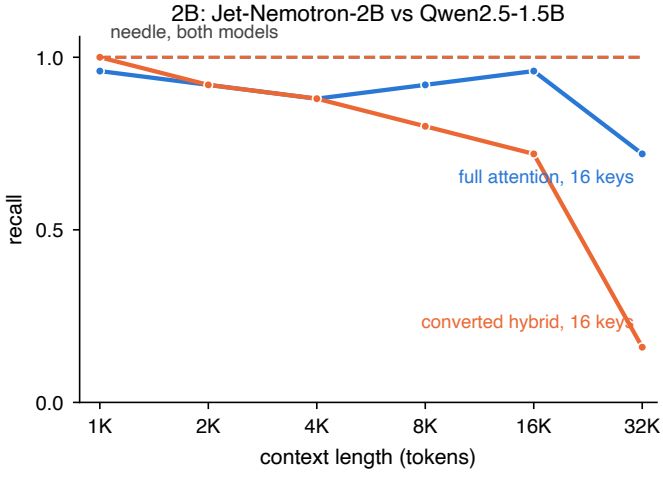
2B: Jet-Nemotron-2B vs Qwen2.5-1.5B
needle, both models
recall
1.0
0.5
0.0
full attention, 16 keys
converted hybrid, 16 keys
1K
2K
4K
8K
16K
32K
context length (tokens)

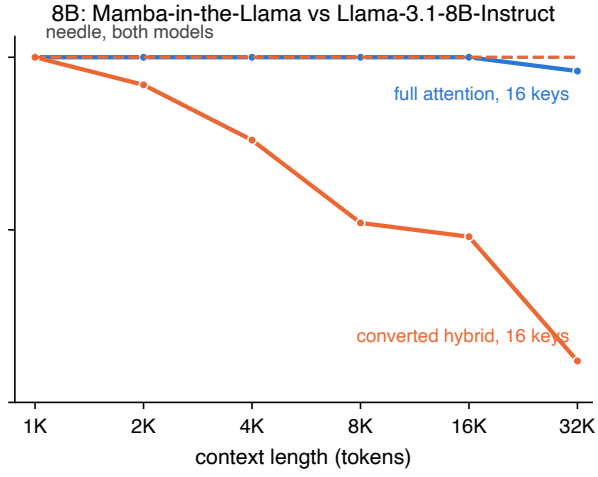
8B: Mamba-in-the-Llama vs Llama-3.1-8B-Instruct
needle, both models
full attention, 16 keys
converted hybrid, 16 keys
1K
2K
4K
8K
16K
32K
context length (tokens)